

MODULE 5: NATURAL LANGUAGE PROCESSING

Contents

1. Learning.....	1
2. Supervised Learning.....	2
3. Classification.....	3
4. Regression.....	4
5. Unsupervised Learning.....	5
6. Clustering in Unsupervised Learning	6
7. Reinforcement Learning	6
8. Natural Language Processing.....	7
9. Artificial Neural Networks (ANNs)	9
10. Syntactic Processing.....	10
11. Semantic Analysis	11
12. Discourse and Pragmatic Processing.....	11
13. Representing and Using Domain Knowledge	12
14. Expert System Shells.....	12
15. Knowledge Acquisition	12

Learning

- 1. Inductive Learning:** Inductive learning involves generalizing from specific examples to broader concepts. It focuses on identifying patterns in a dataset to create a model that can predict outcomes for unseen instances. **Applications:** Classifying emails as spam or not spam, predicting house prices based on historical data.
 - Utilizes training data to infer general rules or relationships.
 - Examples include supervised learning algorithms like decision trees, support vector machines, and neural networks.
- 2. Learning Decision Trees:** Decision tree learning is a method used for classification and regression tasks. It represents decisions and their possible consequences as a tree structure, where each node represents a feature or attribute, branches represent decision rules, and leaves represent outcomes. **Applications:** Customer segmentation, medical diagnosis, credit scoring.
 - Constructs a tree based on the feature that provides the highest information gain or minimizes impurity (e.g., Gini impurity or entropy).
 - Easy to interpret and visualize.
- 3. Explanation-Based Learning (EBL)** Explanation-based learning is a form of learning that uses background knowledge and explanations to generalize from specific instances. It focuses on understanding why a particular example is correct and using that understanding to improve generalization. **Applications:** Natural language understanding, learning concepts in educational software.

- The learner generates explanations for observed behaviour and extracts general rules from these explanations.
 - Often involves deductive reasoning and requires a rich knowledge base.
4. **Learning Using Relevance Information:** This approach involves using additional information to improve learning efficiency and accuracy by focusing on relevant features or instances. It emphasizes filtering out irrelevant data that may hinder the learning process. **Applications:** Document retrieval systems, personalized recommendations, and medical diagnosis where only relevant patient history is considered
 - Can use techniques like feature selection or instance selection to identify and retain only the most relevant data points.
 - Helps in reducing noise and dimensionality in the dataset.
 5. **Neural Net Learning:** Neural network learning is inspired by the structure and functioning of the human brain. It involves the use of interconnected nodes (neurons) organized in layers to learn complex patterns and relationships in data. **Applications:** Image recognition, natural language processing, speech recognition, and game playing (e.g., AlphaGo).
 - Capable of handling large datasets and learning non-linear relationships.
 - Involves techniques like backpropagation for training the network by adjusting weights based on errors.
 6. **Genetic Learning:** Genetic learning is inspired by the principles of natural selection and genetics. It uses evolutionary algorithms to optimize solutions by simulating the process of evolution through selection, crossover, and mutation. **Applications:** Feature selection, optimization problems in engineering, game strategies, and evolving neural network architectures.
 - Population-based approach, where multiple potential solutions (individuals) evolve over generations.
 - Effective in solving complex optimization problems where traditional methods may struggle.

Supervised Learning

Supervised learning is a type of machine learning where an algorithm is trained on labelled data, meaning the input data is paired with the correct output. The goal is for the model to learn from this data and make accurate predictions on new, unseen data. **Key Concepts:**

- **Labelled Data:** The training data consists of input-output pairs (e.g., images labelled as "cat" or "dog").
- **Training:** The model learns by mapping inputs to the correct outputs through iterative learning.
- **Loss Function:** Measures how far the model's predictions are from the true outputs. The model minimizes this function during training.
- **Prediction:** After training, the model can predict outputs for new inputs.

Steps in Supervised Learning

1. **Data Collection:** Gather a labelled dataset with input-output pairs (e.g., images with labels, sales data with outcomes).
2. **Data Preprocessing:** Clean and prepare the data by handling missing values, normalizing features, and splitting into training and test sets.
3. **Feature Selection/Extraction:** Identify the most relevant features or create new features that will help the model make accurate predictions.

4. **Model Selection:** Choose an appropriate supervised learning algorithm (e.g., linear regression, decision tree, SVM, etc.) based on the problem type (classification or regression).
5. **Model Training:** Feed the training data into the chosen algorithm, allowing it to learn patterns from the data by adjusting its parameters.
6. **Model Evaluation:** Test the model using a separate test set to evaluate its performance using metrics such as accuracy, precision, recall, or mean squared error (MSE).
7. **Hyperparameter Tuning:** Optimize the model's hyperparameters (e.g., learning rate, number of layers) to improve its performance.
8. **Prediction:** Use the trained model to make predictions on new, unseen data.
9. **Model Deployment:** Implement the model into a system where it can be used to make real-time predictions or decisions.
10. **Model Monitoring and Maintenance:** Continuously monitor the model's performance over time and retrain it if necessary to maintain accuracy.


Challenges in Supervised Learning:

1. **Data Labelling:** Requires a large amount of labelled data, which can be expensive and time-consuming to obtain.
2. **Overfitting:** The model may perform well on training data but poorly on new, unseen data if it memorizes rather than generalizes.
3. **Imbalanced Data:** When one class is overrepresented in the data, the model may become biased towards that class.
4. **Scalability:** Training models on very large datasets can be computationally expensive and time-consuming.
5. **Noise in Data:** Errors or inconsistencies in the labelled data can lead to incorrect learning and poor model performance.
6. **Limited Generalization:** A model trained on one dataset may not perform well on a different dataset if the data distributions are different.
7. **Feature Engineering:** Identifying the right features can be difficult and may require domain expertise.

Classification


Classification is a type of supervised learning where the goal is to predict a categorical label or class for a given input based on the patterns learned from labelled data. The input is assigned to one of several predefined categories.

1. **Logistic Regression:** Statistical models that use logistic functions to model binary dependent variables. Outputs probabilities and can be used for binary classification tasks. Use information: Email search spam.




2. **Decision Trees:** Decision tree model. Each node represents a feature (or attribute), each branch represents a rule, and each leaf represents a result. Reference: Customer segmentation.


3. **k-Nearest Neighbours (k-NN):** Nonparametric classification technique that assigns labels to points based on their k-nearest training examples within a particular space. Makes use of a distance metric, usually the Euclidean distance, to determine neighbours. Using data: Handwritten digit recognition.



4. **Support Vector Machines (SVM):** The educational supervision model allows it to have the best visibility in separating different classes. Good performance at high altitudes. Data usage: Image classification.



5. **Random Forest:** An ensemble method that combines multiple decision trees to improve accuracy and reduce overfitting. It aggregates the predictions of several trees to make the final decision.
6. **Neural Networks:** Especially with deep learning, neural networks are used for complex classification tasks like image recognition, where input-output relationships are highly non-linear.




Regression

Regression is a type of supervised learning task that aims to predict a continuous numerical value based on input features. Unlike classification, where the output is discrete, regression deals with estimating relationships among variables and forecasting continuous outcomes.


- **Dependent Variable:** The output variable that you are trying to predict (e.g., house prices).
- **Independent Variables:** The input features used to make predictions (e.g., size of the house, location, number of bedrooms).
- **Loss Function:** Measures the difference between the predicted values and the actual values (commonly Mean Squared Error).
- **Training:** The process of adjusting the model parameters to minimize the loss function using historical data.

1. **Linear Regression:** Determine the relationship of dependent with one or more independent variables by fitting a linear equation to observed data. Use data: Predict home values based on various characteristics of home values.



2. **Multiple Linear Regression:** An extension of linear regression that uses multiple independent variables to predict a dependent variable. **Use Cases:** Predicting a person's weight based on height, age, and diet.


3. **Polynomial Regression:** By specifying polynomial terms of input features, it can achieve linear regression. It allows modelling non-linear relationships between these input features. **Possible use:** model growth rates in biology or finance.




4. **Ridge Regression (L2 Regularization):** A linear regression model that adds a regularization term (penalty) to the cost function to reduce overfitting by minimizing the impact of less significant features

5. **Lasso Regression:** Like Ridge regression, different choices can be made by using the L1 constant and reducing some coefficients to zero. **Use data:** Select the most useful features in the data with multiple estimates.

6. **5. Decision Trees Regression:** The algorithm splits the data into subsets depending on feature value and computes the mean value of the target variable for each subset to construct a model tree. **Usage data:** Based on usage patterns and environment to predict device lifetime.



7. **Random Forest Regression:** It is an ensemble approach that involves large numbers of decision trees to multiply the effectiveness of any single prediction. Each tree is trained on a random subset of the data, and the final prediction is the mean prediction of the individual trees.



Unsupervised Learning

Unsupervised learning is a type of machine learning where the model is trained on data without labelled outcomes. The goal is to discover hidden patterns or intrinsic structures within the input data. Unlike supervised learning, there is no predefined output, and the model learns from the raw data itself. **Key Concepts:**

- **No Labels:** Data used in unsupervised learning does not have associated target outputs or labels.
- **Data Exploration:** The focus is on exploring the data and finding relationships, clusters, or patterns.
- **Dimensionality Reduction:** Techniques used to reduce the number of features while retaining essential information.

Applications of Unsupervised Learning

1. **Market Segmentation:** Grouping customers based on purchasing behaviour to tailor marketing strategies.
2. **Anomaly Detection:** Identifying fraudulent transactions or network intrusions by detecting outliers in data.
3. **Recommendation Systems:** Suggesting products or content to users based on similar preferences found in data.
4. **Customer Profiling:** Understanding customer behaviour patterns to improve user experience.
5. **Image Compression:** Reducing the size of images while retaining important features using dimensionality reduction techniques.
6. **Document Clustering:** Grouping similar documents together for information retrieval or topic modelling.

Challenges in Unsupervised Learning

1. **Evaluation:** Lack of ground truth makes it difficult to evaluate the quality of results.
2. **Interpretability:** Understanding and interpreting the patterns or clusters found can be complex.
3. **Scalability:** Some algorithms may struggle to scale effectively with large datasets.
4. **Choosing the Right Algorithm:** Selecting the appropriate technique for a specific task can be challenging without labelled data.
5. **Data Quality:** The effectiveness of unsupervised learning heavily depends on the quality and preprocessing of the input data.

Clustering in Unsupervised Learning

Clustering is a type of unsupervised learning that involves grouping a set of objects (data points) into clusters based on their similarities. The objective is to ensure that objects within the same cluster are more like each other than to those in other clusters. Clustering helps in identifying natural groupings in data, making it easier to analyse and interpret. **Key Concepts:**

- **Distance Metrics:** Measures used to quantify the similarity or dissimilarity between data points. Common metrics include Euclidean distance, Manhattan distance, and cosine similarity.
 - **Centroids:** The central point of a cluster, often used in algorithms like K-Means to represent the cluster.
 - **Inertia:** A measure of how tightly the clusters are packed together, often used to assess the quality of the clustering.
1. **K-Means Clustering:** Partitions the data into K clusters by assigning each data point to the nearest centroid and updating the centroids based on the mean of the points in each cluster. **Use Cases:** Customer segmentation, image compression. **Limitations:** Requires the number of clusters (K) to be specified in advance and is sensitive to outliers.
 2. **Hierarchical Clustering:** Builds a hierarchy of clusters using either a bottom-up approach (agglomerative) or a top-down approach (divisive). The result is often represented in a dendrogram. **Use Cases:** Taxonomy classification, social network analysis. **Limitations:** Computationally expensive for large datasets.

Reinforcement Learning

Reinforcement Learning (RL) is a type of machine learning where an agent learns to make decisions by interacting with an environment. The agent takes actions in this environment to maximize a cumulative

reward over time. Unlike supervised learning, where the model is trained on labeled data, in reinforcement learning, the agent learns from the consequences of its actions.

- **Agent:** The learner or decision-maker that interacts with the environment.
- **Environment:** The setting in which the agent operates. The agent observes the state of the environment and takes actions.
- **State:** A representation of the current situation of the agent within the environment.
- **Action:** The choices available to the agent to interact with the environment.
- **Reward:** A feedback signal received by the agent after taking an action in a particular state, indicating the success of that action.
- **Policy:** A strategy that the agent employs to determine which actions to take in different states. It can be deterministic or stochastic.
- **Value Function:** A function that estimates the expected cumulative reward that can be obtained from a given state or state-action pair.

Applications of Reinforcement Learning:

1. **Game Playing:** RL has been successfully applied in games like chess, Go, and video games (e.g., AlphaGo and OpenAI's Dota 2 agent).
2. **Robotics:** Training robots to perform tasks such as walking, grasping objects, or navigating environments.
3. **Autonomous Vehicles:** Making real-time decisions based on sensor inputs to navigate and respond to dynamic environments.
4. **Healthcare:** Personalizing treatment plans for patients based on their responses to different interventions.
5. **Recommendation Systems:** Adapting recommendations based on user interactions and feedback.

Challenges in Reinforcement Learning:

1. **Exploration vs. Exploitation:** Balancing the need to explore new actions (to discover their rewards) with the need to exploit known actions that yield high rewards.
2. **Sample Efficiency:** Many RL algorithms require a large number of interactions with the environment to learn effectively, which can be impractical in real-world applications.
3. **Sparse Rewards:** In some environments, the agent may receive very few rewards, making it difficult to learn effective policies.
4. **Generalization:** Ensuring that learned policies can generalize well to unseen states or environments.
5. **Stability and Convergence:** Ensuring that the learning process is stable and converges to optimal policies can be challenging, especially in complex environments.

Natural Language Processing

Natural Language Processing (NLP) is a subfield of Artificial Intelligence (AI) focused on enabling machines to understand, interpret, and generate human language in a way that is both meaningful and useful. NLP combines linguistics, computer science, and AI to bridge the gap between human communication and machine understanding. **Key Components of NLP:**

1. **Tokenization:** Breaking down a sentence into words or smaller parts called tokens.
2. **Lemmatization and Stemming:** Reducing words to their root form or base meaning.
 - a. **Stemming:** Trimming words to their base form (e.g., "running" to "run").
 - b. **Lemmatization:** Reducing words to their dictionary form (e.g., "better" to "good").
3. **Stop Words Removal:** Removing common words like 'is', 'the', 'and' which add little meaning.

4. **Named Entity Recognition (NER):** Identifying proper nouns like names of people, places, organizations, etc.
5. **Part of Speech (POS) Tagging:** Identifying the grammatical role of each word (e.g., noun, verb, adjective).
6. **Sentiment Analysis:** Determining the emotional tone behind a piece of text (positive, negative, or neutral).
7. **Text Classification:** Assigning predefined categories to text (e.g., spam detection in emails).
8. **Word Embeddings:** Representing words in a numerical format, often used for understanding semantic similarities between words (e.g., Word2Vec, GloVe).

Applications of NLP:

1. **Chatbots and Virtual Assistants:** Provide automated, real-time customer support and perform tasks based on voice commands.
2. **Machine Translation:** Automatically translate languages (e.g., Google Translate), facilitating cross-lingual communication.
3. **Sentiment Analysis:** Analyse text data to understand opinions in social media, reviews, or feedback.
4. **Text Summarization:** Generate concise summaries of long documents for faster information consumption.
5. **Speech Recognition:** Convert spoken language into text for transcription or interaction with systems.
6. **Spam Detection:** Classify and filter out unwanted or harmful emails using text analysis.
7. **Healthcare:** Assist in medical document analysis, summarization of reports, and extracting useful patient information.

Advantages of NLP:

1. **Automation of Routine Tasks:** Reduces human effort in text processing, like customer service and translation.
2. **Enhanced Customer Experience:** Provides faster, personalized support via chatbots and virtual assistants.
3. **Better Insights from Data:** Helps extract meaningful information from large volumes of unstructured text data.
4. **Improved Decision-Making:** Facilitates better business decisions through sentiment analysis and text mining.
5. **Accessibility:** Enables interaction with devices through natural language, improving accessibility for users with disabilities.
6. **Scalability:** NLP solutions can handle vast amounts of text data efficiently, making them scalable across industries.
7. **Cross-Language Communication:** Helps bridge communication gaps between speakers of different languages through real-time translation.

Challenges in NLP:

1. **Ambiguity:** Words or sentences can have multiple meanings, making it hard for machines to interpret correctly.
2. **Context Understanding:** Grasping the broader context or intent behind words is complex for NLP systems.
3. **Sarcasm and Irony:** Difficult to detect subtle expressions like sarcasm or humor, leading to incorrect interpretations.

4. **Low-Resource Languages:** Limited data and resources for many languages make it harder to build NLP models for them.
5. **Data Privacy:** Processing sensitive personal data like emails or conversations raises privacy concerns.
6. **Idioms and Slang:** Unique cultural expressions are tough for machines to understand accurately.
7. **Real-Time Processing:** Ensuring efficient, accurate language understanding in real-time applications is challenging.

Artificial Neural Networks (ANNs)

Definition: Artificial Neural Networks (ANNs) are computational models inspired by the biological neural networks found in the human brain. They consist of interconnected groups of nodes (neurons) that work together to process input data and generate output. ANNs are particularly effective for tasks involving pattern recognition, classification, regression, and data modelling.

- **Neurons:** The fundamental building blocks of an ANN, analogous to biological neurons. Each neuron receives input, processes it, and produces an output.
- **Layers:**
 - **Input Layer:** The first layer that receives the input data.
 - **Hidden Layers:** Intermediate layers where neurons process inputs. ANNs can have one or more hidden layers, enabling them to learn complex patterns.
 - **Output Layer:** The final layer that produces the output of the network.
- **Weights and Biases:** Each connection between neurons has an associated weight, which determines the importance of that connection. Biases are additional parameters added to neurons to adjust the output.
- **Activation Functions:** Functions applied to the output of neurons to introduce non-linearity into the model. Common activation functions include:
 - **Sigmoid:** Outputs values between 0 and 1.
 - **Tanh (Hyperbolic Tangent):** Outputs values between -1 and 1.
 - **ReLU (Rectified Linear Unit):** Outputs the input directly if positive; otherwise, it outputs zero.

How ANNs Work:

1. **Forward Propagation:** Data is fed into the input layer, and it propagates through the network layer by layer. Each neuron calculates its output based on the inputs, weights, biases, and activation function.
2. **Loss Function:** The difference between the predicted output and the actual output (ground truth) is quantified using a loss function. Common loss functions include Mean Squared Error for regression tasks and Cross-Entropy Loss for classification tasks.
3. **Backpropagation:** The network adjusts its weights and biases based on the error computed by the loss function. This is done using an optimization algorithm (e.g., Stochastic Gradient Descent) to minimize the loss.
4. **Training:** The process of repeatedly feeding input data through the network and adjusting weights through backpropagation continues until the model converges, meaning it achieves satisfactory performance on the training data.

Types of Artificial Neural Networks:

1. **Feedforward Neural Networks:** Information moves in one direction—from input to output—without loops. Used for simple tasks like regression and classification.
2. **Convolutional Neural Networks (CNNs):** Designed specifically for processing grid-like data, such as images. They use convolutional layers to detect patterns and features, making them ideal for image recognition tasks.
3. **Recurrent Neural Networks (RNNs):** Designed for sequential data, RNNs have connections that loop back, allowing them to maintain memory of previous inputs. Commonly used for tasks like natural language processing and time-series forecasting.
4. **Long Short-Term Memory (LSTM) Networks:** A special type of RNN that effectively handles long-range dependencies by using memory cells and gating mechanisms to control the flow of information.
5. **Generative Adversarial Networks (GANs):** Consist of two networks (generator and discriminator) that compete against each other. The generator creates fake data, while the discriminator evaluates whether the data is real or fake. Used for image generation, style transfer, and more.

Applications of Artificial Neural Networks:

1. **Image Recognition:** Identifying objects, faces, or scenes in images (e.g., facial recognition systems).
2. **Natural Language Processing:** Tasks such as sentiment analysis, translation, and chatbots.
3. **Speech Recognition:** Converting spoken language into text and vice versa.
4. **Medical Diagnosis:** Analysing medical images and predicting diseases based on patient data.
5. **Financial Predictions:** Stock market forecasting and credit scoring.
6. **Game Playing:** Developing intelligent agents that can play complex games (e.g., AlphaGo).

Advantages of ANNs:

1. **Flexibility:** Capable of learning complex relationships and patterns from data.
2. **Generalization:** Can generalize learned patterns to unseen data, making them robust in various applications.
3. **Scalability:** Can handle large datasets and improve performance with more data.

Challenges of ANNs:

1. **Data Requirements:** Require large amounts of labelled data for effective training.
2. **Computational Resources:** Can be resource-intensive, requiring significant computational power and memory.
3. **Overfitting:** Risk of learning noise and outliers in the training data, which can lead to poor performance on new data.
4. **Interpretability:** Often considered "black boxes," making it difficult to understand the reasoning behind their predictions.

Syntactic Processing

Syntactic processing involves analysing the grammatical structure of sentences to ensure they conform to the rules of a specific language grammar. It determines how words are organized to form meaningful phrases or sentences. **Parsing Techniques:**

1. **Top-Down Parsing:**

- Starts with the root of a parse tree and expands it using grammar rules until the input sentence is matched.
- **Advantage:** Focuses directly on deriving the input sentence.
- **Example:** Starting from "S" (Sentence) and expanding to "NP + VP" (Noun Phrase + Verb Phrase).

2. **Bottom-Up Parsing:**

- Begins with the input words and attempts to combine them into larger components until the root of the parse tree is reached.
- **Advantage:** Ensures input is valid before constructing higher-level grammar.
- **Example:** Combining "the cat" into a Noun Phrase (NP).

Grammar Types:

1. **Context-Free Grammar (CFG):**

- A set of production rules where the left-hand side of each rule contains a single non-terminal symbol.
- **Example:** $S \rightarrow NP VP$ (Sentence $\rightarrow$ Noun Phrase + Verb Phrase).

2. **Dependency Grammar:**

- Focuses on the relationships (dependencies) between words in a sentence.
- **Example:** In "The cat sleeps," "sleeps" is the root, and "cat" depends on "sleeps."

Semantic Analysis

Semantic analysis aims to understand the **meaning** of words, phrases, and sentences. It ensures that the derived meaning is contextually accurate. **Techniques:**

1. **Semantic Roles:** Assign roles to sentence components based on their function (e.g., agent, action, object). **Example:** In "John ate an apple," "John" is the agent, "ate" is the action, and "apple" is the object.
2. **Word Sense Disambiguation (WSD):** Determines the correct meaning of a word based on context. **Example:** For "bank," WSD decides whether it refers to a "riverbank" or a "financial institution" depending on surrounding words.

Representations:

1. **Semantic Networks:** Graph-based structures representing relationships between concepts. **Example:** "Cat $\rightarrow$ isA $\rightarrow$ Animal."
2. **Knowledge Graphs:** A more advanced representation connecting entities and their relationships. **Example:** "Barack Obama $\rightarrow$ wasPresidentOf $\rightarrow$ USA."

Discourse and Pragmatic Processing

Discourse: Focuses on understanding the relationships between sentences or phrases in a text. **Key Tasks:**

- **Coherence:** Ensuring logical flow between sentences. **Example:** Linking "John bought a car. He drove it home" (linking "He" to "John").
- **Coreference Resolution:** Identifying when two expressions refer to the same entity. **Example:** "The dog barked. It was hungry" (linking "It" to "The dog").

Pragmatic Processing: Interprets meaning based on the **context** and **user intent**, going beyond the literal meaning. **Example:** Sentence: "Can you pass the salt?" **Pragmatic meaning:** A request, not a literal question about ability. Includes understanding sarcasm, politeness, and implied meanings.

Representing and Using Domain Knowledge

1. **Knowledge Representation:** Techniques for organizing and structuring domain-specific knowledge to be used by expert systems:
 - **Rules:** Represented as IF-THEN statements for decision-making. **Example:** IF temperature > 100°F THEN alert = "High Fever."
 - **Frames:** Structured data representations using attributes and values. **Example:** Frame for "Patient" might include attributes like "Name," "Age," and "Symptoms."
 - **Semantic Networks:** Graph-based structures linking concepts through relationships. **Example:** "Flu → causes → Fever."
1. **Rule-Based Systems:**
 - Decision-making systems based on predefined rules.
 - Rules are triggered by specific conditions, leading to a conclusion or action.
 - **Example:** A medical system diagnosing a disease based on symptoms.

Expert System Shells

A software framework for building expert systems without starting from scratch. Provides a generic structure that developers customize for specific domains. **Components:**

- **Knowledge Base:** Stores facts and rules about the domain. **Example:** Disease symptoms and treatments in a medical expert system.
- **Inference Engine:** Applies logical reasoning to the knowledge base to derive conclusions. **Example:** Matches patient symptoms to rules to suggest a diagnosis.
- **User Interface:** Allows users to interact with the system, input data, and view results. **Example:** A graphical interface for entering symptoms and receiving diagnoses.

Knowledge Acquisition

1. **Techniques:** Methods for gathering domain-specific knowledge to populate the knowledge base:
 - **From Domain Experts:** Interviewing or observing experts to extract their knowledge. **Example:** Consulting doctors for medical knowledge.
 - **From Databases:** Using existing data sources to identify patterns and relationships. **Example:** Analyzing hospital records for diagnostic rules.
2. **Tools:** Automated tools for extracting and organizing knowledge:
 - **Machine Learning:** Algorithms learn patterns and rules from data. **Example:** Training a model on patient data to predict diseases.
 - **Data Mining:** Techniques to discover patterns and relationships in large datasets. **Example:** Identifying correlations between symptoms and diseases.