

Module 1.2: Introduction to Hadoop Ecosystem

Contents

Hadoop	1
Hadoop Ecosystem	2
Open-Source Technologies.....	4
Cloud and Big Data.....	5
Cloud Storage	6
Cloud Analytics.....	6
Big Data as a Service (BDaaS)	6
Mobile Business Intelligence (Mobile BI)	7
Mobile Analytics.....	7
Real-Time Reporting	7
Mobile Dashboards	8
Crowdsourcing Analytics.....	8
Data Collection from Large Groups	8
Inter-Firewall Analytics	10
Security Monitoring.....	10
Data Sharing Between Organizations	10
Trans-Firewall Analytics	10
Cross-Enterprise Data Analysis.....	10
Security and Privacy Concerns	11

Hadoop

Hadoop: Hadoop is an open-source framework developed by the Apache Software Foundation for storing and processing large volumes of data across multiple computers in a distributed environment. It is designed to handle structured, semi-structured, and unstructured data efficiently using clusters of commodity hardware.

Need for Hadoop

Traditional database systems face limitations when dealing with massive amounts of data. Hadoop was developed to overcome these challenges.

- **Handling Big Data:** Processes huge volumes of data that traditional systems cannot handle efficiently.
- **Scalability:** New nodes can be added easily as data grows.
- **Cost Effectiveness:** Uses low-cost commodity hardware instead of expensive servers.
- **Fault Tolerance:** Continues operation even if some nodes fail.
- **Distributed Processing:** Processes data in parallel across multiple machines, improving performance.

Hadoop Features

- **Distributed Storage:** Stores data across multiple machines.
- **Distributed Processing:** Processes data simultaneously on different nodes.
- **Large-Scale Data Processing:** Handles terabytes and petabytes of data efficiently.
- **Scalability:** Supports growth from a few nodes to thousands of nodes.
- **High Processing Speed:** Uses parallel processing techniques.

- **Open Source:** Freely available and supported by a large community.
- **Fault Tolerance:** Automatically replicates data to prevent loss.
- **High Availability:** Ensures continuous access to data.
- **Flexibility:** Supports structured, semi-structured, and unstructured data.
- **Cost Efficiency:** Uses inexpensive hardware for storage and processing.

Hadoop Ecosystem

Hadoop Ecosystem: A collection of tools and technologies that work together with Hadoop to store, process, analyze, and manage Big Data.

1. **HDFS (Hadoop Distributed File System):** is the primary storage system of Hadoop. It is designed to store very large files across multiple machines in a distributed environment. Instead of storing a file on a single computer, HDFS divides the file into blocks and stores them on different nodes in a cluster.

Working: When a file is uploaded to HDFS, it is broken into smaller blocks (typically 128 MB or 256 MB). These blocks are distributed among different DataNodes. To ensure reliability, each block is replicated and stored on multiple nodes. If one node fails, the data can still be accessed from another node.

Features

- Distributed storage of large files.
- Automatic data replication.
- Fault tolerance.
- High scalability.
- High throughput access to data.


Advantages

- Reliable data storage.

- Supports petabytes of data.
- Fault recovery through replication.
- Cost-effective storage solution.

Applications

- Data warehousing.
- Log storage.
- Social media data storage.
- Large-scale analytics.



2. **MapReduce:** is the data processing framework of Hadoop. It processes large datasets in parallel by dividing the work among multiple machines and combining the results to generate final output.

Working: MapReduce consists of two phases:

- **Map Phase:** The input data is divided into smaller chunks and processed independently. Each mapper generates intermediate key-value pairs.
- **Reduce Phase:** The reducer collects the intermediate key-value pairs, groups similar keys together, and produces the final output.

Features

- Parallel processing.
- Distributed computation.
- Automatic fault handling.
- Scalable processing framework.

Advantages

- Processes massive datasets efficiently.

- Reduces execution time.
- Supports distributed computing.
- Improves system performance.

Applications

- Log analysis.
- Data mining.
- Search engine indexing.
- Business analytics.

3. **YARN (Yet Another Resource Negotiator):** is the resource management and job scheduling framework of Hadoop. It manages cluster resources and allocates them to different applications running in the Hadoop environment.

Working: When a job is submitted, YARN determines which resources are required and allocates CPU, memory, and storage resources accordingly. It continuously monitors running applications and ensures efficient resource utilization.

Main Components of YARN

- **Resource Manager:** Manages and allocates cluster resources.
- **Node Manager:** Runs on each node and monitors resource usage.
- **Application Master:** Manages the execution of a particular application.
- **Containers:** Units where application tasks are executed.

Features

- Efficient resource management.
- Multi-user support.
- Job scheduling.
- Cluster monitoring.

Advantages

- Better cluster utilization.
- Improved scalability.

- Supports multiple processing frameworks.
- Efficient workload management.

Applications

- Resource allocation.
- Cluster management.
- Job scheduling.

4. **Hive:** is an open-source data warehouse system built on Hadoop that allows users to query large datasets using HiveQL, a language similar to SQL. It enables users who are familiar with SQL to analyze Big Data without writing complex MapReduce programs.

Working: Hive converts HiveQL queries into MapReduce, Tez, or Spark jobs and executes them on Hadoop.

Features

- SQL-like query language.
- Data summarization.
- Query optimization.
- Data warehousing support.

Advantages

- Easy to learn for SQL users.
- Simplifies data analysis.
- Supports large datasets.

- Reduces programming complexity.

Applications

- Business intelligence.
- Report generation.
- Data analysis.
- Data warehousing.

Example

SELECT * FROM Employees;

5. **Pig:** is an open-source data processing platform used for analysing large datasets on Hadoop. It uses a scripting language called Pig Latin. Pig simplifies complex data processing tasks that would otherwise require lengthy MapReduce code.

Working: Users write Pig Latin scripts which are automatically converted into MapReduce jobs and executed on Hadoop.

Features

- High-level scripting language.
- Supports structured and unstructured data.
- Reduces coding effort.
- Easy data transformation.

Advantages

- Faster development.
- Less coding than MapReduce.

- Suitable for ETL operations.
- Easy debugging.

Applications

- Data transformation.
- Data cleansing.
- ETL processes.
- Data analysis.

Example

A = LOAD 'employee.txt';

6. **HBase:** is a distributed NoSQL database built on top of HDFS. It is designed for storing huge amounts of structured and semi-structured data while providing real-time read and write access.

Features

- Column-oriented database.
- Distributed architecture.
- Real-time access.
- High scalability.

Advantages

- Fast data retrieval.
-

- Handles billions of records.
- Fault tolerant.
- Supports large-scale applications.

Applications

- Social media applications.
- Online transaction systems.
- Real-time analytics.

7. **Apache Spark:** is an open-source distributed computing framework that performs high-speed data processing in memory instead of repeatedly reading from disk. It is considered much faster than traditional MapReduce processing.

Working: Spark loads data into memory and performs computations directly in RAM, reducing disk access and improving processing speed.

Features

- In-memory computing.
- Real-time processing.
- Machine learning support.
- Stream processing.
- Graph processing.

Advantages

- Faster than MapReduce.

- Supports batch and real-time analytics.
- Easy integration with Hadoop.
- Supports advanced analytics.

Applications

- Machine learning.
- Data analytics.
- Real-time processing.
- Recommendation systems.

Open-Source Technologies

Open Source Technology refers to software whose source code is freely available to the public. Users can view, modify, distribute, and improve the software according to their requirements. Open-source software is usually developed collaboratively by a community of developers and organizations.

- Source code is publicly accessible.
- Users can modify and customize the software.
- No licensing fees in most cases.
- Community-driven development.
- Continuous improvement through collaboration.

Benefits of Open Source Technologies

- 1. Cost Effectiveness:** Open-source software is generally available free of cost, reducing software licensing expenses.
- 2. Flexibility and Customization:** Users can modify the source code according to their specific requirements.
- 3. Transparency:** The source code is publicly available, allowing users to inspect and verify the software.
- 4. Community Support:** Large developer communities continuously contribute updates, bug fixes, and improvements.
- 5. Reliability:** Continuous testing and contributions make open-source software stable and dependable.
- 6. Security:** Vulnerabilities can be identified and fixed quickly due to public code review.
- 7. Scalability:** Open-source technologies can easily handle increasing amounts of data and workloads.
- 8. Vendor Independence:** Organizations are not locked into a single vendor and have complete control over the software.
- 9. Faster Innovation:** Continuous contributions from developers worldwide help introduce new features rapidly.
- 10. Easy Integration:** Most open-source tools can be integrated easily with other technologies and platforms.

Examples of Open-Source Technologies: Hadoop, Hive, Pig, Apache Spark

Comparison of Open-Source Big Data Technologies

Technology	Primary Purpose
Hadoop	Distributed storage and processing
Hive	SQL-based data analysis
Pig	Data processing using scripts
Spark	Fast in-memory data processing
Cassandra	Distributed NoSQL database

Cloud and Big Data

Cloud Computing is the delivery of computing resources such as servers, storage, databases, networking, and software over the Internet on a pay-as-you-use basis.

- On-demand access to resources.
- Scalability and flexibility.
- Resource sharing.
- Remote accessibility.
- Pay-per-use pricing model.

Service Models

- **IaaS (Infrastructure as a Service):** Provides virtual machines, storage, and networking resources.
- **PaaS (Platform as a Service):** Provides a platform for developing and deploying applications.
- **SaaS (Software as a Service):** Provides software applications over the Internet.

Cloud Storage

Cloud Storage is a service that allows data to be stored, managed, and accessed over the Internet instead of local storage devices. **Features**

- Large storage capacity.
- Automatic backup and recovery.
- Data accessibility from anywhere.
- High availability and reliability.
- Easy scalability.
- Microsoft OneDrive
- Amazon S3

Importance in Big Data: Cloud storage enables organizations to store massive volumes of structured, semi-structured, and unstructured data without investing in expensive hardware.

Examples

- Google Drive
- Dropbox

Cloud Analytics

Cloud Analytics refers to the use of cloud-based tools and services to process, analyze, and visualize large volumes of data. **Functions**

- Data collection.
- Data processing.
- Data analysis.
- Reporting and visualization.
- Predictive analytics.
- Real-time analytics.
- Reduced infrastructure costs.
- Easy integration with Big Data tools.

Examples

- Google BigQuery
- Amazon Redshift
- Microsoft Azure Synapse Analytics

Advantages

- Faster data processing.

Benefits of Cloud for Big Data

1. **Scalability:** Cloud platforms can easily increase or decrease resources based on data requirements.
2. **Cost Efficiency:** Organizations pay only for the resources they use, reducing infrastructure costs.
3. **High Storage Capacity:** Cloud environments can store enormous amounts of Big Data.
4. **Faster Processing:** Powerful cloud resources enable rapid processing of large datasets.
5. **Flexibility:** Resources can be allocated dynamically according to workload demands.
6. **Accessibility:** Data and applications can be accessed from anywhere through the Internet.
7. **Reliability:** Cloud providers offer backup, disaster recovery, and fault-tolerant systems.
8. **Real-Time Analytics:** Cloud platforms support real-time data processing and decision-making.
9. **Easy Deployment:** Big Data applications can be deployed quickly without purchasing hardware.
10. **Enhanced Collaboration:** Multiple users can access and analyze data simultaneously.

Big Data as a Service (BDaaS)

Big Data as a Service (BDaaS) is a cloud-based service model that provides Big Data storage, processing, analytics, and management capabilities over the Internet.

Working: Organizations upload data to a cloud platform where the service provider manages infrastructure, storage, processing, and analytics tools.

Features

- Cloud-based Big Data processing.
- On-demand analytics services.
- Reduced infrastructure management.
- Pay-as-you-go pricing.
- Scalable resources.

Advantages

1. **Lower Cost:** Eliminates the need for expensive hardware and maintenance.
2. **Scalability:** Resources can be increased or decreased as needed.
3. **Faster Implementation:** Big Data solutions can be deployed quickly.
4. **Reduced Complexity:** Service providers handle infrastructure management.

5. **Improved Accessibility:** Data and analytics tools are accessible from anywhere.

Examples of BDaaS Platforms

- Amazon Web Services (AWS)
- Google Cloud Platform (GCP)
- Microsoft Azure
- IBM Cloud

Relationship Between Cloud and Big Data

Cloud computing provides the infrastructure, storage, and processing power required for Big Data applications. Big Data generates enormous volumes of information, while cloud platforms offer scalable and cost-effective solutions for storing, processing, and analyzing that data. Together, they enable organizations to extract valuable insights and make data-driven decisions efficiently.

Mobile Business Intelligence (Mobile BI)

Mobile Business Intelligence (Mobile BI) refers to the delivery of Business Intelligence (BI) tools, reports, dashboards, and analytics on mobile devices such as smartphones and tablets, enabling users to access business information anytime and anywhere.

- Access to business data through mobile devices.
- Real-time monitoring of business performance.
- Interactive reports and dashboards.
- Supports decision-making on the move.

Mobile Analytics

Mobile Analytics is the process of collecting, measuring, and analyzing business data through mobile applications and devices to gain insights and support decision-making.

Uses

- Tracking sales performance.
- Monitoring customer behavior.
- Analyzing business trends.
- Measuring application usage.

Benefits

- Quick access to data insights.
- Better understanding of customer activities.
- Faster business decisions.

Real-Time Reporting

Real-Time Reporting provides up-to-date business information immediately after data is generated, allowing organizations to respond quickly to changing situations.

Features

- Instant data updates.
- Live monitoring of operations.
- Immediate identification of issues.
- Supports rapid decision-making.

Examples

- Live sales reports.
- Stock market updates.
- Website traffic monitoring.
- Fraud detection alerts.

Mobile Dashboards

Mobile Dashboards are visual interfaces designed for mobile devices that display key business metrics, reports, and performance indicators in an easy-to-understand format.

Components

- Charts and graphs.
- Performance indicators (KPIs).
- Reports and summaries.
- Alerts and notifications.

Benefits

- Easy visualization of data.
- Quick performance monitoring.
- Improved accessibility.
- User-friendly interface

Advantages of Mobile BI

- 1. Anywhere Access:** Users can access business information from any location using mobile devices.
- 2. Real-Time Decision Making:** Provides instant access to current business data for faster decisions.
- 3. Increased Productivity:** Employees can monitor and manage business activities while away from the office.
- 4. Improved Collaboration:** Teams can share reports and insights quickly through mobile platforms.
- 5. Better Customer Service:** Real-time information helps organizations respond to customer needs more effectively.
- 6. Faster Reporting:** Business reports can be viewed instantly without waiting for desktop access.
- 7. Enhanced Monitoring:** Managers can continuously track business performance and key metrics.
- 8. Cost Efficiency:** Reduces dependence on physical infrastructure and office-based reporting systems.

Crowdsourcing Analytics

Crowdsourcing Analytics is the process of collecting, analyzing, and utilizing data contributed by a large number of people through online platforms, mobile applications, social media, and other digital channels. **Key Features**

- Data is gathered from a large group of individuals.
- Uses internet-based platforms for participation.
- Provides diverse and real-time information.
- Supports large-scale data collection.

Data Collection from Large Groups

Data Collection from Large Groups involves gathering information, opinions, feedback, or observations from thousands or millions of participants.

Methods

- Online surveys.
- Social media contributions.
- Mobile applications.
- User reviews and ratings.
- Crowdsourcing platforms.

Importance

- Generates large datasets quickly.
- Provides diverse viewpoints.
- Reduces data collection costs.
- Improves data availability.

Crowdsourced Data Sources

- **Social Media:** Platforms like Facebook, X (Twitter), and Instagram generate user-created content.
- **Online Reviews:** Customer reviews on e-commerce and service websites provide valuable insights.
- **Mobile Applications:** Apps collect user feedback, location data, and activity information.
- **Crowdsourcing Platforms:** Websites specifically designed for collecting information from large groups.
- **Community Forums:** Discussion boards and online communities contribute opinions and experiences.
- **Wikipedia:** Information is created and maintained by a large community of contributors.

Applications of Crowdsourcing Analytics

1. **Market Research:** Organizations analyze customer opinions and preferences to improve products and services.
2. **Sentiment Analysis:** Businesses study social media content to understand public attitudes toward brands and products.
3. **Traffic Monitoring:** Navigation applications collect user-generated location data to provide real-time traffic updates.
4. **Disaster Management:** Information from citizens helps authorities respond quickly during emergencies.
5. **Product Development:** Customer suggestions and feedback help organizations develop better products.
6. **Healthcare Research:** Large groups contribute health-related information for medical studies and disease tracking.
7. **Content Creation:** Crowdsourced information helps build knowledge bases and online communities.

Benefits

1. **Large Data Volume:** Enables collection of massive amounts of data from numerous participants.
2. **Cost Effective:** Reduces the cost of traditional data collection methods.
3. **Real-Time Information:** Data can be collected and analyzed quickly.
4. **Diverse Perspectives:** Information comes from people with different backgrounds and experiences.
5. **Scalability:** Easily gathers data from a large population.
6. **Faster Decision Making:** Organizations gain timely insights for better decisions.

Challenges

1. **Data Quality Issues:** User-contributed data may be inaccurate or incomplete.
2. **Data Reliability:** Not all contributors provide trustworthy information.
3. **Privacy Concerns:** Collection of user data may raise privacy and security issues.
4. **Data Validation:** Large datasets require verification before analysis.
5. **Bias in Contributions:** Data may not represent all population groups equally.
6. **Security Risks:** Crowdsourced platforms can be vulnerable to false or malicious information.

Inter-Firewall Analytics

Inter-Firewall Analytics/Analytics Across Multiple Networks refers to the process of collecting and analyzing data from multiple networks protected by different firewalls to gain a broader view of operations, performance, and security. **Importance**

- Provides a unified view of distributed networks.
- Helps detect network-wide issues.
- Improves coordination across multiple locations.
- Supports enterprise-wide decision making.

Security Monitoring

Security Monitoring involves continuously analyzing network traffic, logs, and security events across multiple firewall-protected networks to identify threats and suspicious activities.

Benefits

- Early detection of cyber attacks.
- Continuous monitoring of network activities.
- Improved incident response.
- Enhanced network security.

Examples

- Intrusion detection.
- Malware monitoring.
- Unauthorized access detection.
- Threat intelligence analysis

Data Sharing Between Organizations

Data Sharing Between Organizations allows multiple organizations to exchange selected data securely for collaborative analysis while maintaining network protection through firewalls.

Applications

- Banking and financial fraud detection.
- Healthcare research collaborations.
- Supply chain management.
- Government security operations.

Benefits

- Better insights through combined data.
- Improved collaboration.
- Enhanced decision-making.
- More effective threat detection.

Trans-Firewall Analytics

Trans-Firewall Analytics/Analytics Beyond Organizational Boundaries refers to analyzing data that originates from different organizations, departments, or external sources beyond a single firewall-protected environment.

Importance

- Provides broader business insights.
- Supports large-scale analytics.
- Enables collaborative research and innovation.
- Improves strategic planning.

Cross-Enterprise Data Analysis

Cross-Enterprise Data Analysis involves combining and analyzing data from multiple organizations to identify patterns, trends, and opportunities that may not be visible within a single organization.

Applications

- Industry benchmarking.
- Fraud detection across institutions.
- Market research.

- Healthcare analytics.
- Supply chain optimization.

Benefits

- More comprehensive analysis.

- Improved predictive capabilities.
- Better risk assessment.
- Enhanced business intelligence.

Security and Privacy Concerns

Security and Privacy Concerns arise when data is shared and analyzed across organizational boundaries, requiring strict measures to protect sensitive information. **Major Concerns**

- 1. Data Privacy:** Personal and confidential information must be protected from unauthorized access.
- 2. Data Security:** Shared data must be secured during storage and transmission.
- 3. Regulatory Compliance:** Organizations must comply with data protection laws and regulations.
- 4. Access Control:** Only authorized users should access shared data.

- 5. Data Breaches:** Shared environments may increase the risk of cyber attacks and information leakage.

Solutions

- Data encryption.
- Authentication and authorization.
- Secure communication channels.
- Data anonymization.
- Regular security audits.

Difference Between Inter-Firewall and Trans-Firewall Analytics

Inter-Firewall Analytics	Trans-Firewall Analytics
Operates across multiple firewall-protected networks.	Operates beyond organizational boundaries.
Focuses on network-level monitoring and analysis.	Focuses on cross-enterprise data analysis.
Usually within related organizations or departments.	Often involves multiple independent organizations.
Primarily used for security and operational monitoring.	Primarily used for collaboration and large-scale analytics.